Parallel and Distributed Computing with MATLAB

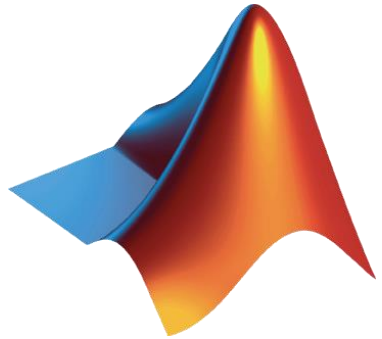

Agenda

- Utilizing multiple cores on a desktop computer
- Scaling up to cluster and cloud resources
- Tackling data-intensive problems on desktops and clusters
- Accelerating applications with NVIDIA GPUs
- Summary and resources

Steps for Improving Performance

- First get code working
- Speed up code with core MATLAB
- Include compiled languages and additional hardware

Webinar: Optimizing and Accelerating Your MATLAB Code

Programming Parallel Applications

- Built-in multithreading
  - Automatically enabled in MATLAB since R2008a
  - Multiple threads in a single MATLAB computation engine

- Parallel computing using explicit techniques
  - Multiple computation engines controlled by a single session
  - High-level constructs to let you parallelize MATLAB applications
  - Perform MATLAB computations on GPUs

Practical Application of Parallel Computing

- Why parallel computing?
  - Need faster insight on more complex problems with larger datasets
  - Computing infrastructure is broadly available (multicore desktops, GPUs, clusters)

- Why parallel computing with MATLAB
  - Leverage computational power of more hardware
  - Accelerate workflows with minimal to no code changes to your original code
  - Focus on your engineering and research, not the computation

Parallel Computing Paradigm

Multicore Desktops

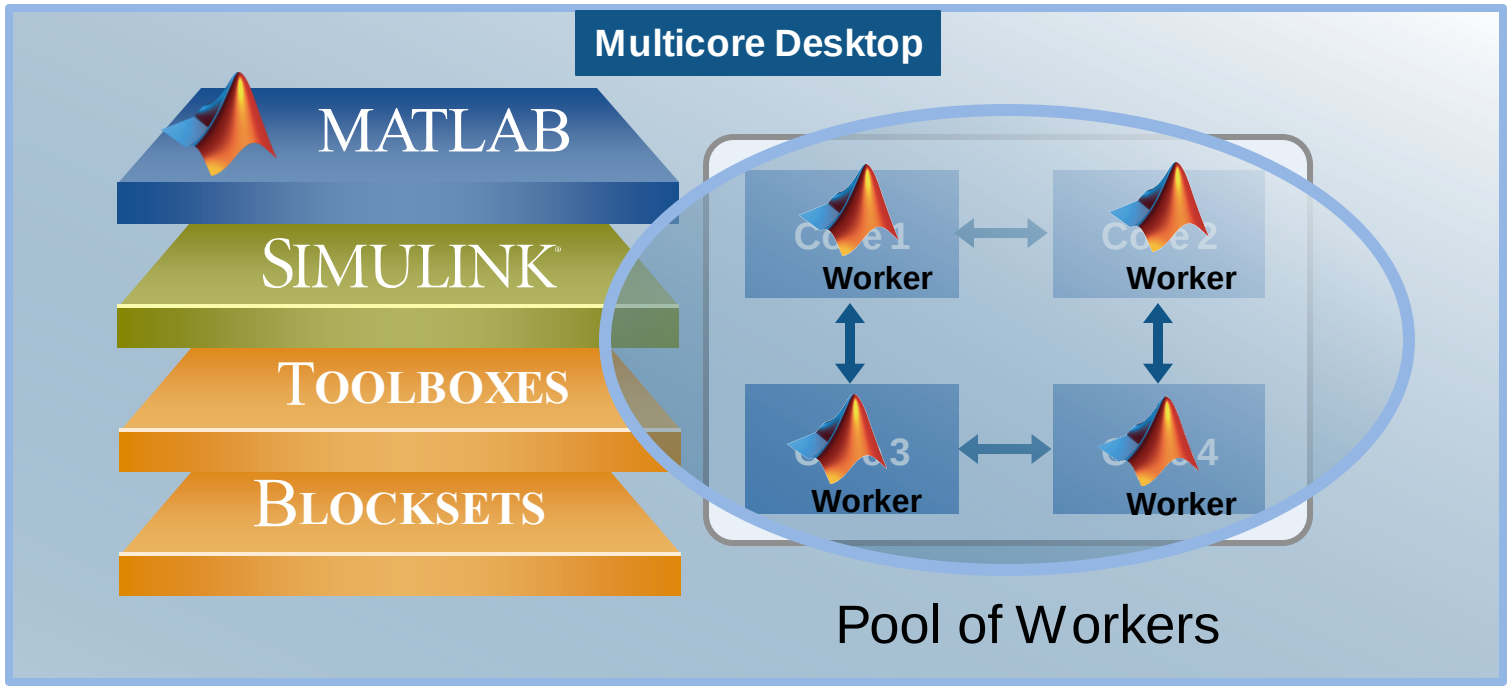

MATLAB multicore

Programming Parallel Applications

- Built in support
  - `..., 'UseParallel', true)`

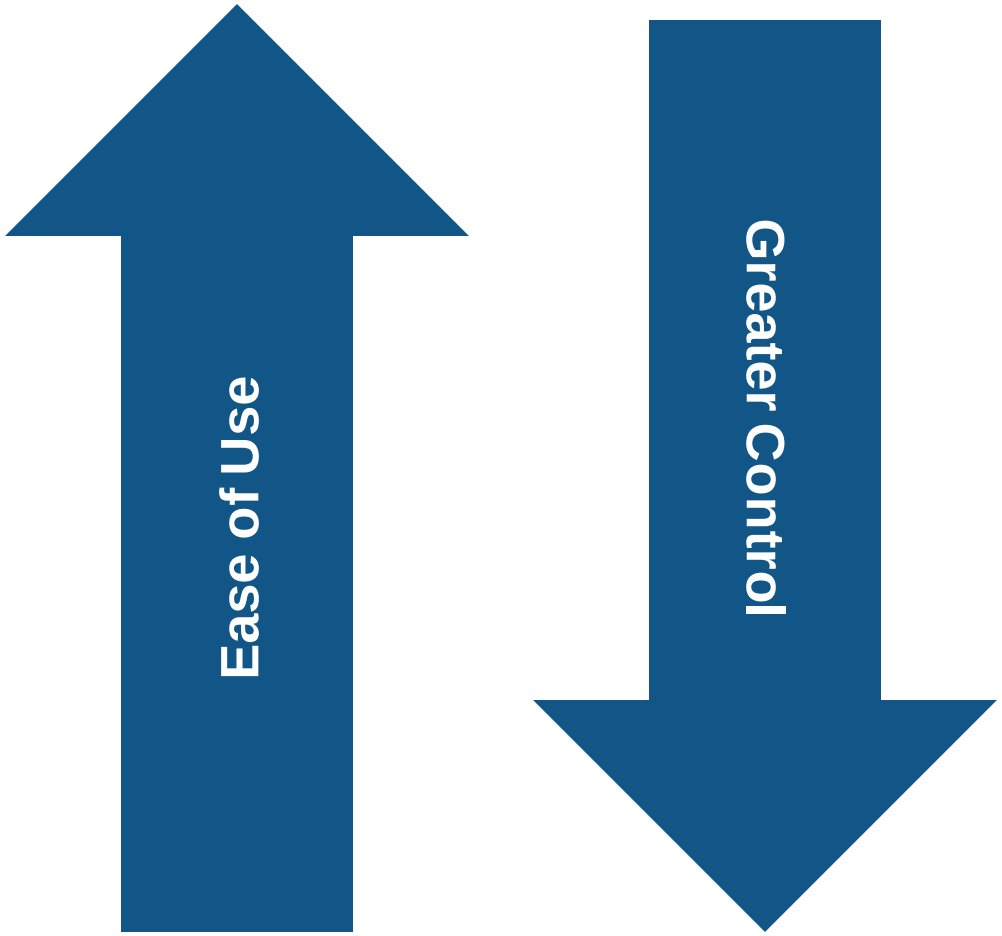

Parallel-enabled Toolboxes

Image Processing

Batch Image Processor, Block Processing, GPU-enabled functions

Statistics and Machine Learning

Resampling Methods, k-Means clustering, GPU-enabled functions

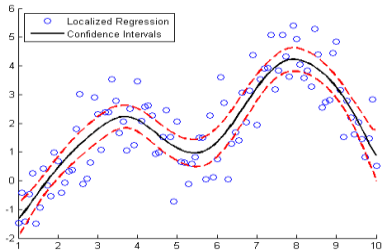

Neural Networks

Deep Learning, Neural Network training and simulation

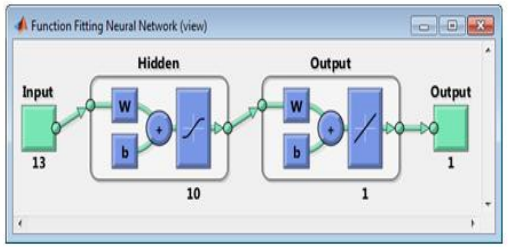

Signal Processing and Communications

GPU-enabled FFT filtering, cross correlation, BER

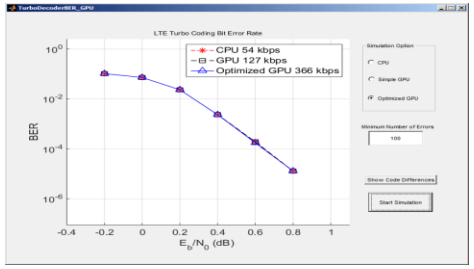

Computer Vision

Parallel-enabled functions in bag-of-words workflow

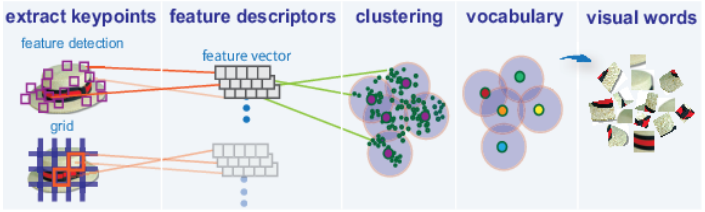

Optimization

Parallel estimation of gradients

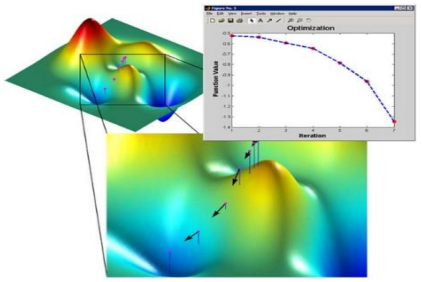

[Other Parallel-enabled Toolboxes]

Programming Parallel Applications

- Built in support
  - `..., 'UseParallel', true)`
- Simple programming constructs
  - `parfor`, `batch`

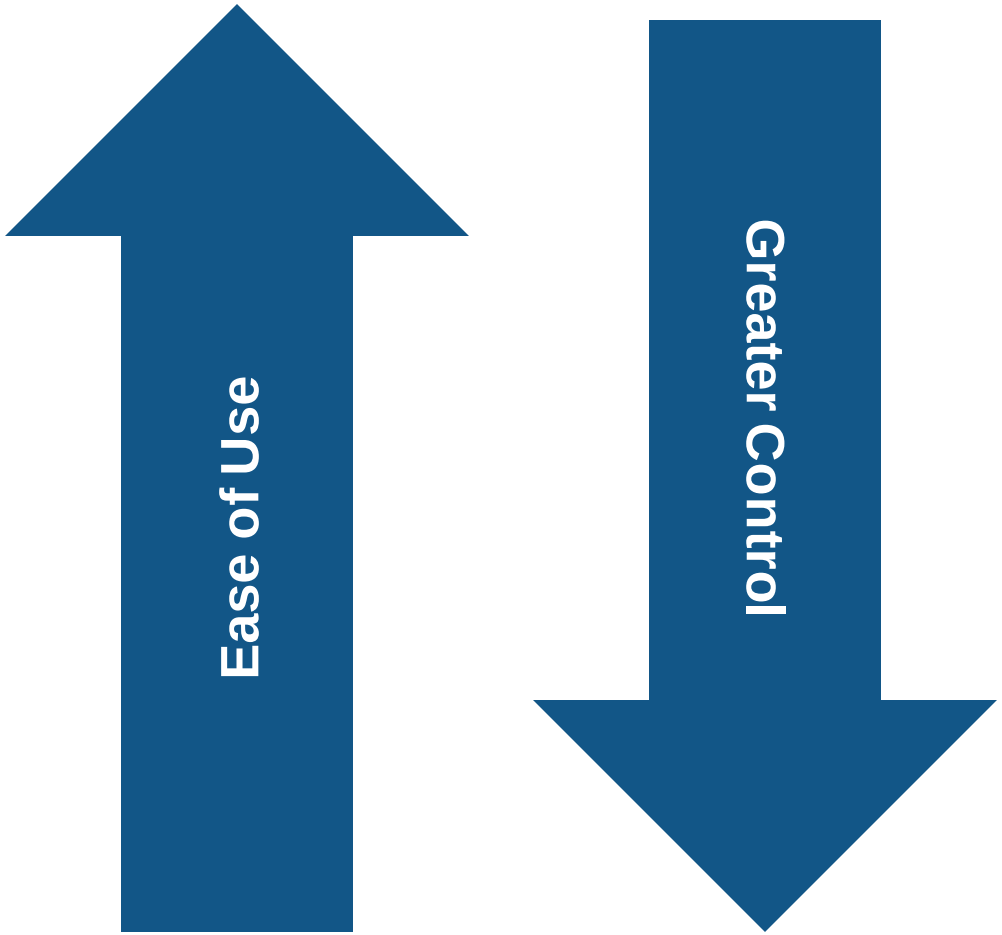

Embarrassingly Parallel: Independent Tasks or Iterations

- No dependencies or communication between tasks
- Examples:
  - Monte Carlo simulations
  - Parameter sweeps
  - Same operation on many files

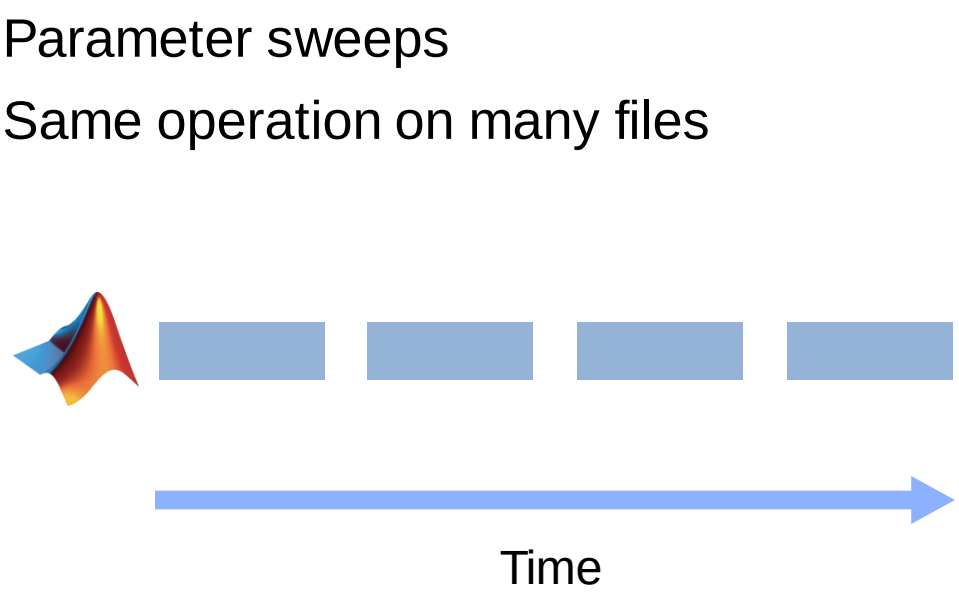

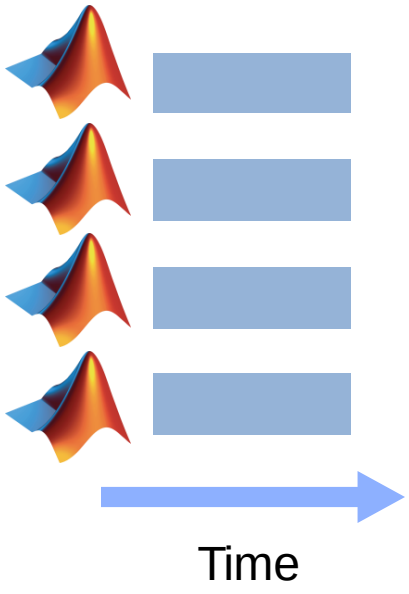

Mechanics of `parfor` Loops

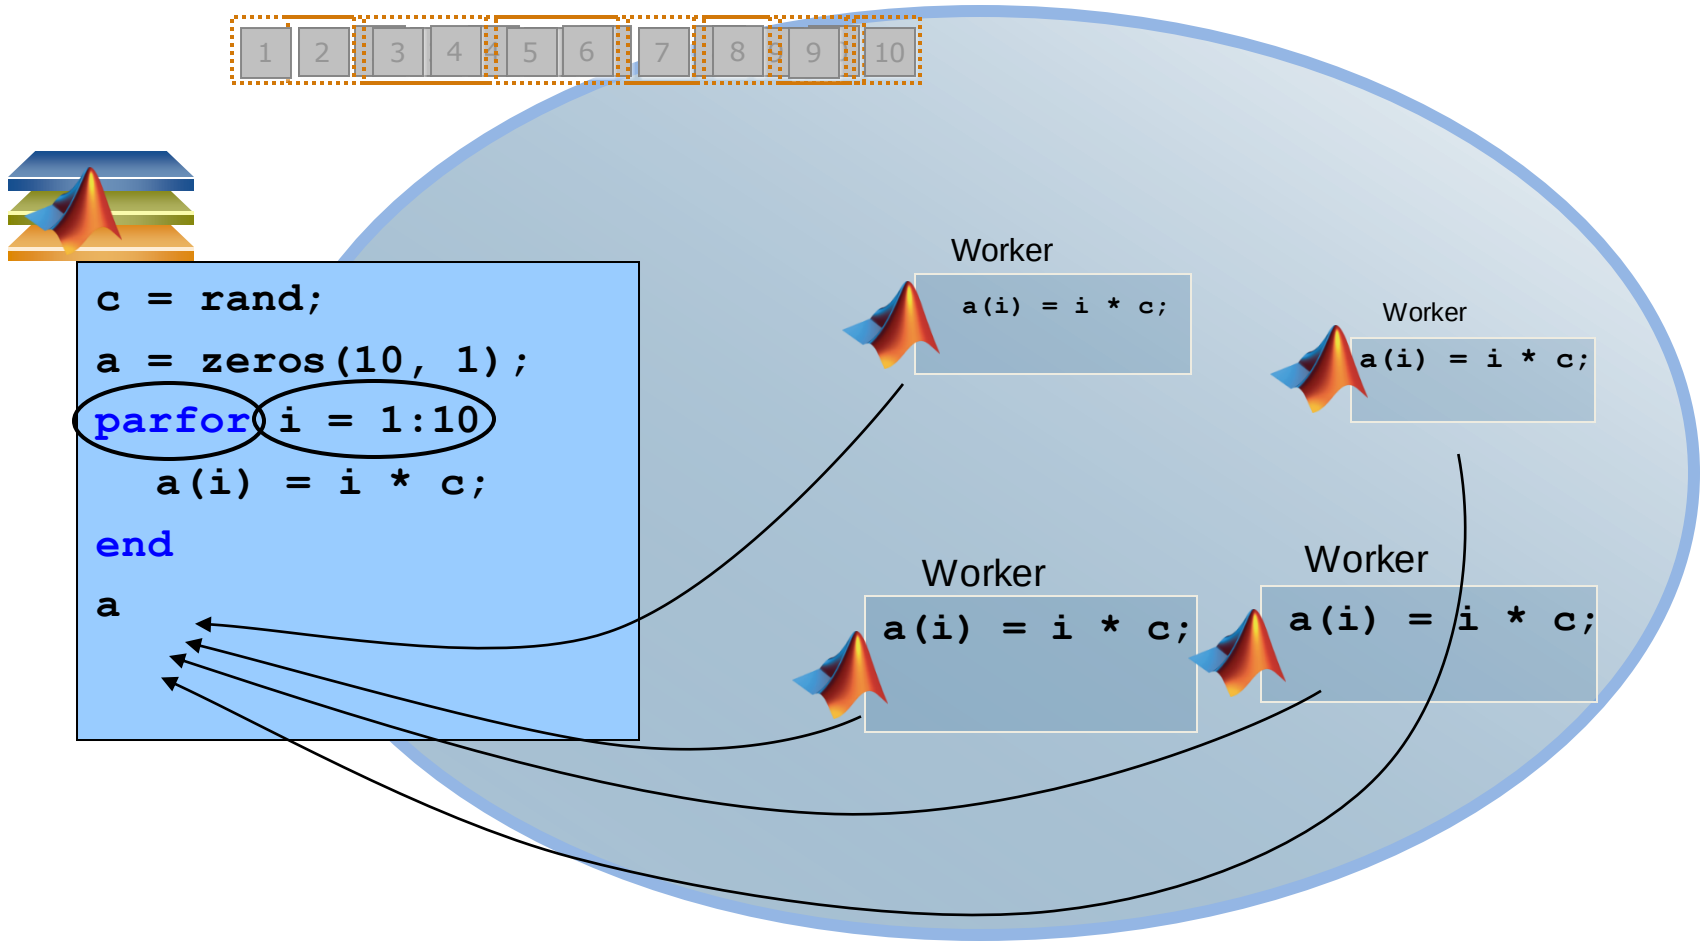

Tips for Leveraging PARFOR

- Consider creating smaller arrays on each worker versus one large array prior to the parfor loop

- Take advantage of `parallel.pool.Constant` to establish variables on pool workers prior to the loop

- Encapsulate blocks as functions when needed

Understanding parfor

Parallel Computing with MATLAB – Beyond PARFOR

Discussed Features

- parallel-enabled toolboxes
- `parfor`

Additional Features

- Batch submission to the cluster
- GPU utilization
- Granular parfor control
- Parallel support for big data
- High memory jobs
- Low level message passing

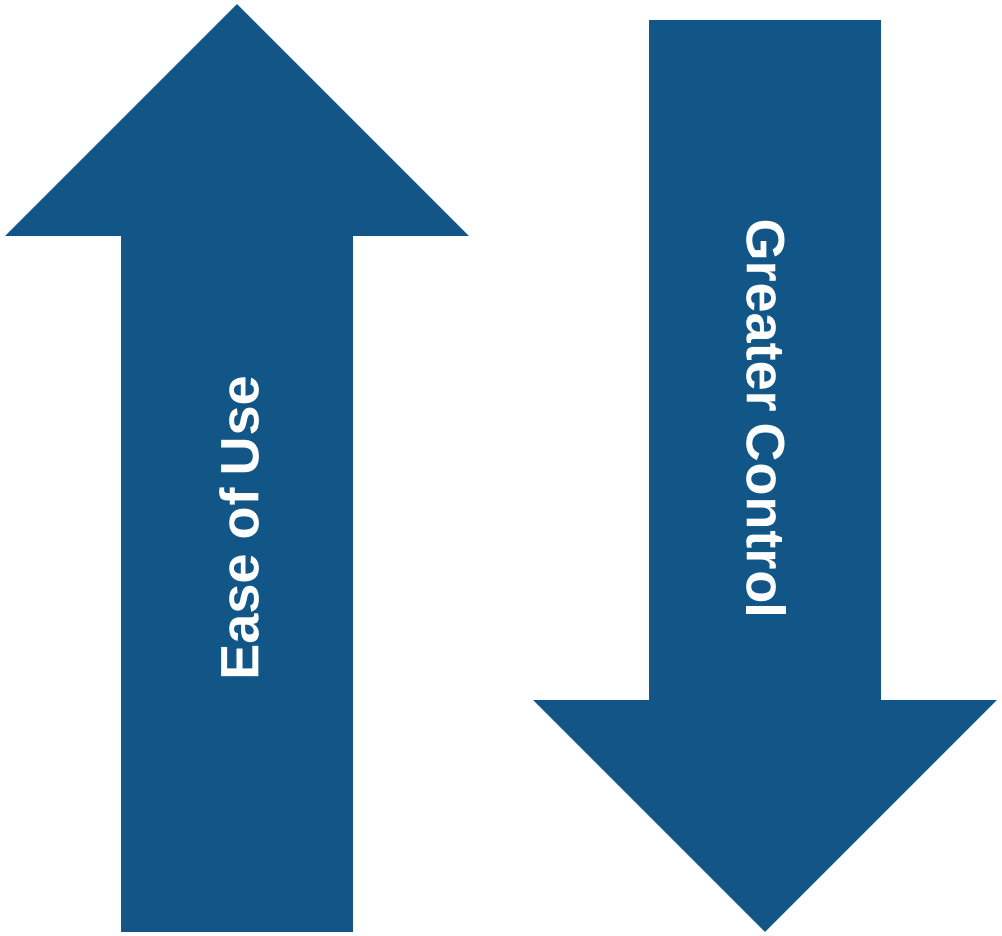

Agenda

- Utilizing multiple cores on a desktop computer
- Scaling up to cluster and cloud resources
- Tackling data-intensive problems on desktops and clusters
- Accelerating applications with NVIDIA GPUs
- Summary and resources

Migrate execution to a cluster environment

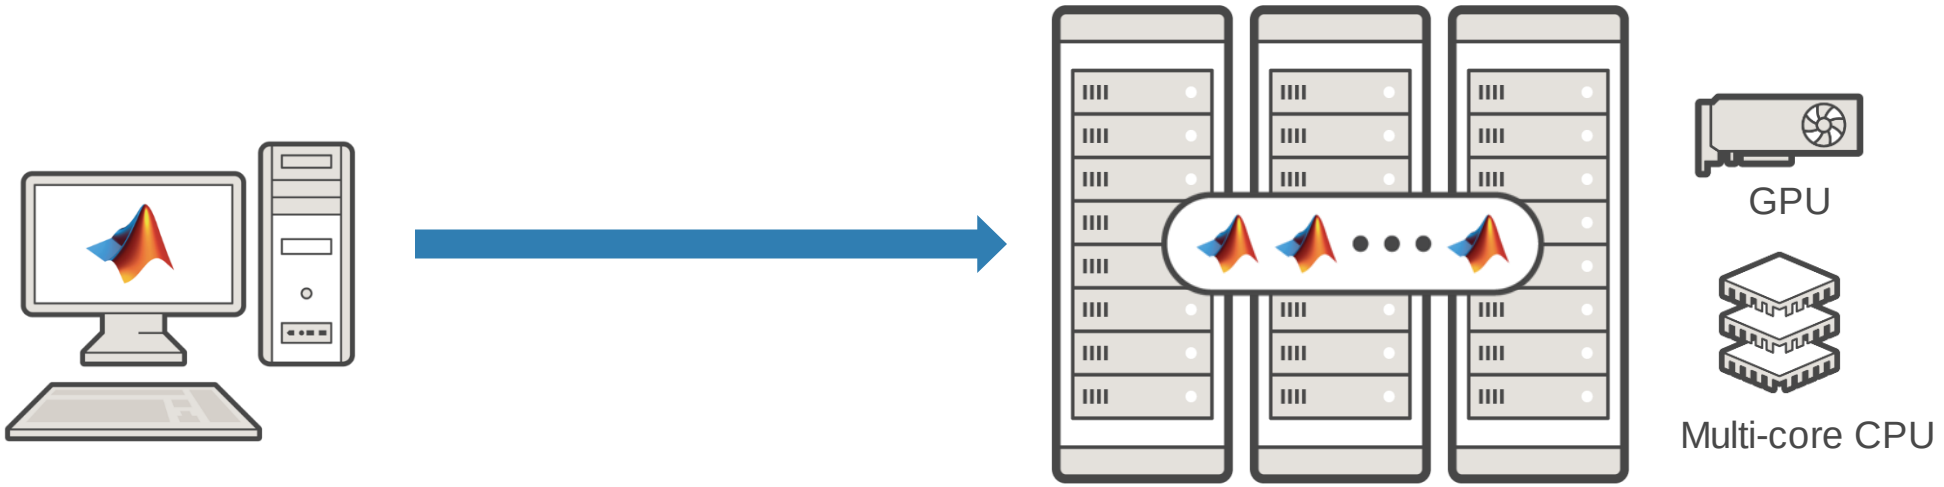

MATLAB

Parallel Computing Toolbox

Prototype on the desktop

MATLAB Distributed Computing Server

Integrate with existing infrastructure

Offloading Serial Computations

- `job = batch(...);`

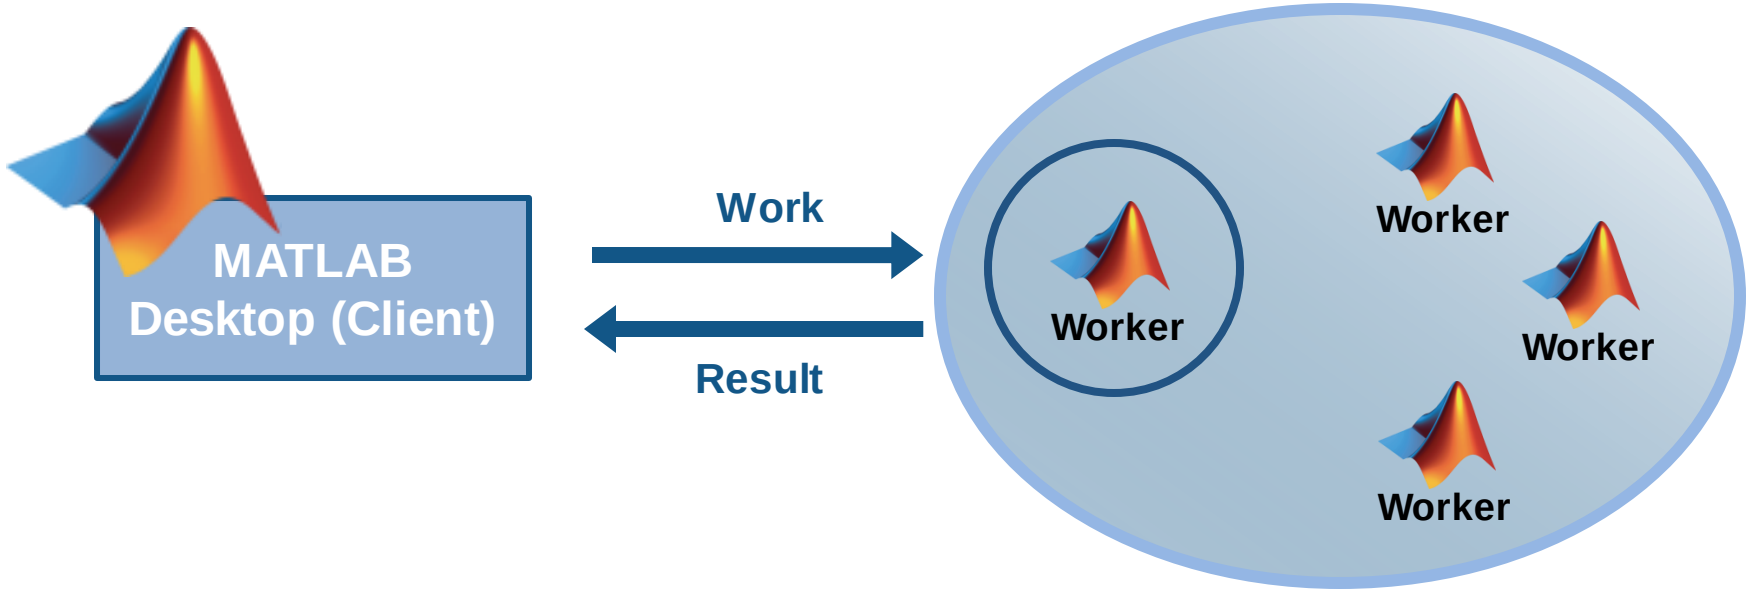

Offloading and Scaling Computations

- `job = batch(... , 'Pool', n);`

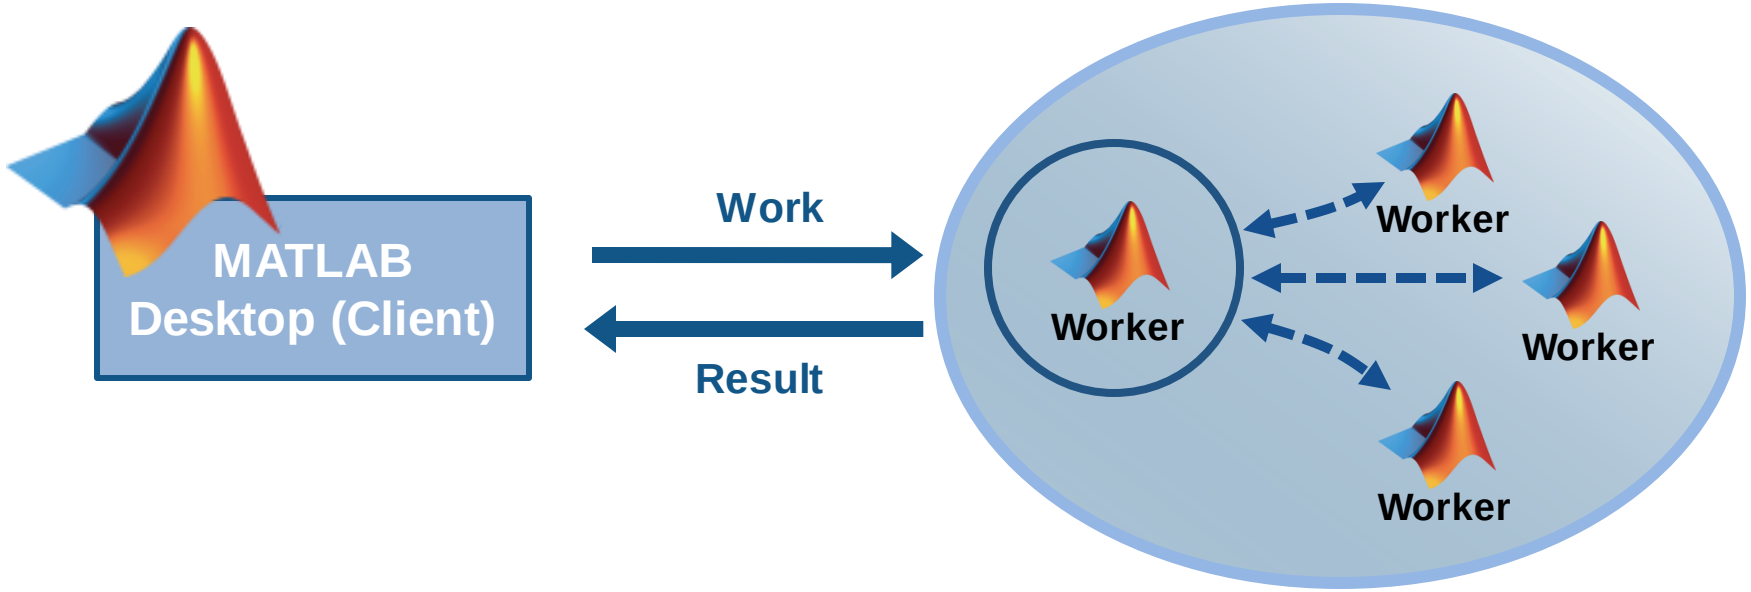

Migrate to Cluster / Cloud

- Use MATLAB Distributed Computing Server
- Change hardware without changing algorithm

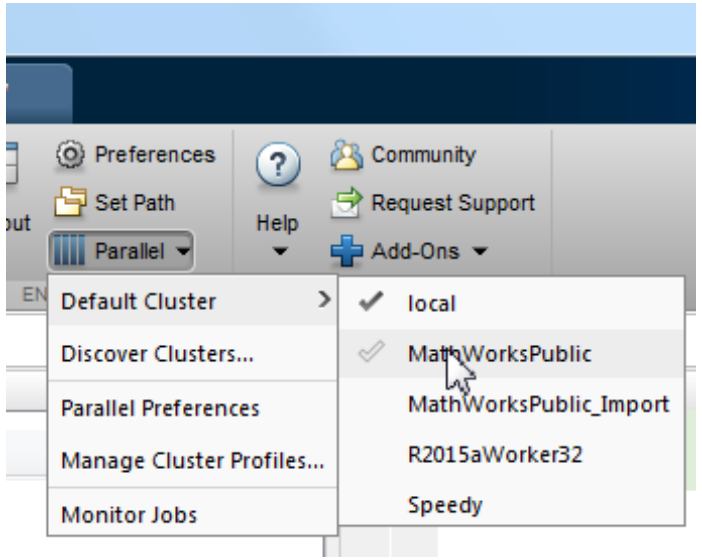

Use MATLAB Distributed Computing Server

1. Prototype code

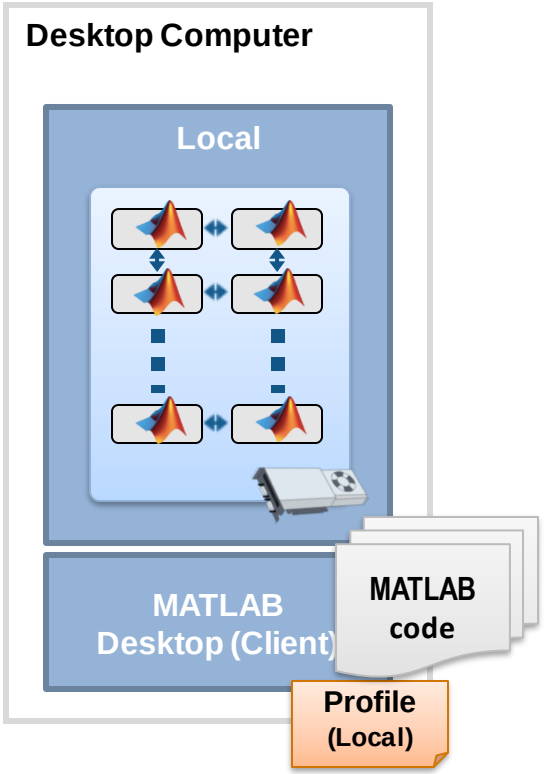

Use MATLAB Distributed Computing Server

1. Prototype code
2. Get access to an enabled cluster

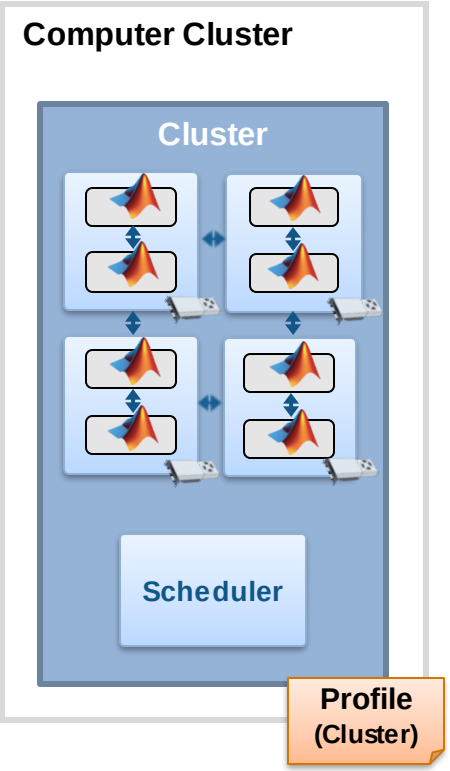

Use MATLAB Distributed Computing Server

1. Prototype code
2. Get access to an enabled cluster
3. Switch cluster profile to run on cluster resources

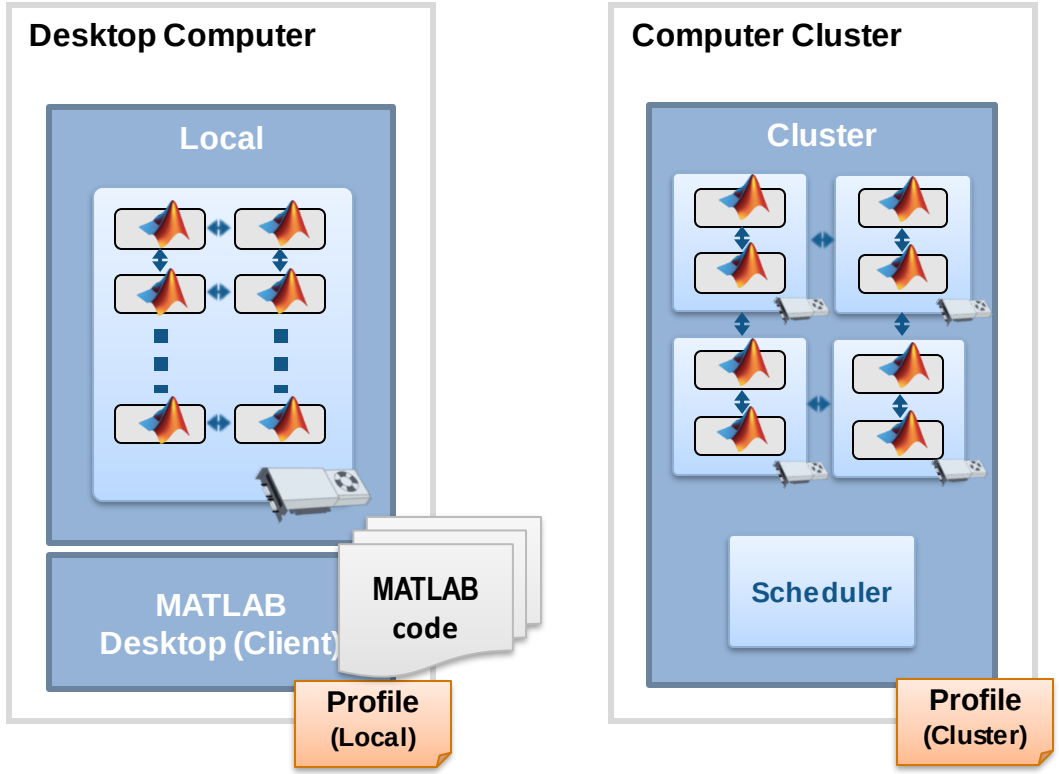

Take Advantage of Cluster Hardware

- Offload computation:
  - Free up desktop
  - Access better computers
- Scale speed-up:
  - Use more cores
  - Go from hours to minutes
- Scale memory:
  - Utilize tall arrays and distributed arrays
  - Solve larger problems without re-coding algorithms

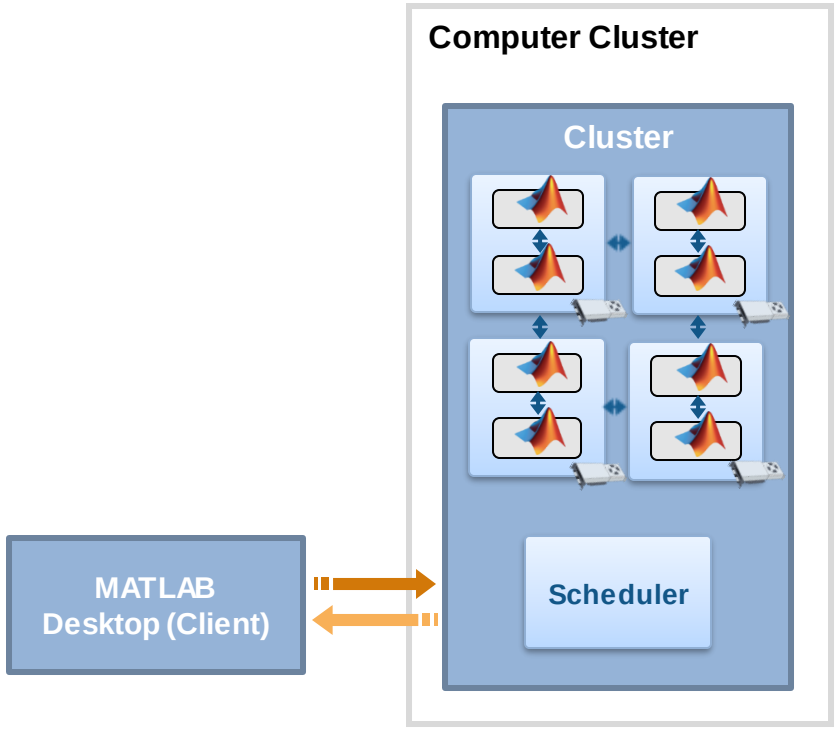

Scale your applications beyond the desktop

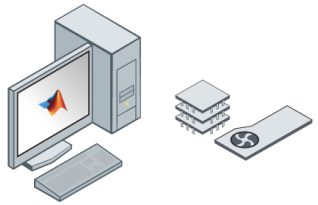

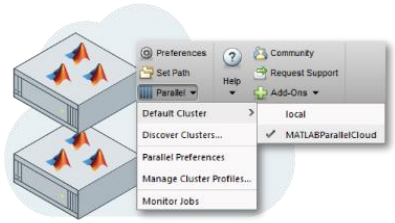

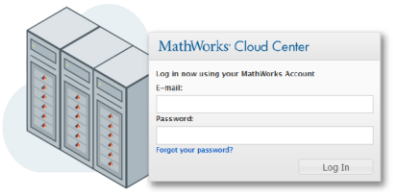

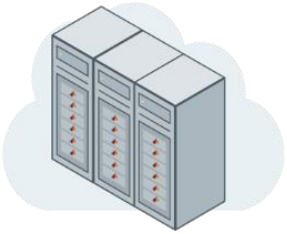

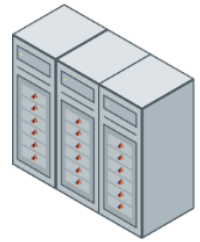

Option	Parallel Computing Toolbox	MATLAB Parallel Cloud	MATLAB Distributed Computing Server *for Amazon EC2*	MATLAB Distributed Computing Server *for Custom Cloud*	MATLAB Distributed Computing Server
Description	Explicit desktop scaling	Single-user, basic scaling to cloud	Scale to EC2 with some customization	Scale to custom cloud	Scale to clusters
Maximum workers	No limit	16	1000	No limit	No limit
Hardware	Desktop	MathWorks Compute Cloud	Amazon	Amazon, Azure, Others	Any
Availability	Worldwide	United States and Canada	United States, Canada and other select countries in Europe	Worldwide	Worldwide

Learn More: Parallel Computing on the Cloud

Agenda

- Utilizing multiple cores on a desktop computer
- Scaling up to cluster and cloud resources
- Tackling data-intensive problems on desktops and clusters
- Accelerating applications with NVIDIA GPUs
- Summary and resources

Tall Arrays

- New data type in MATLAB R2016b
- Applicable when:
  - Data is **columnar** – with **many** rows
  - Overall data size is **too big to fit into memory**
  - Operations are mathematical/statistical in nature
  - Common Actions
    - Data manipulation, math, statistics
    - Summary visualizations
    - Machine learning
- Statistical and machine learning applications
  - Hundreds of functions supported in MATLAB and Statistics and Machine Learning Toolbox

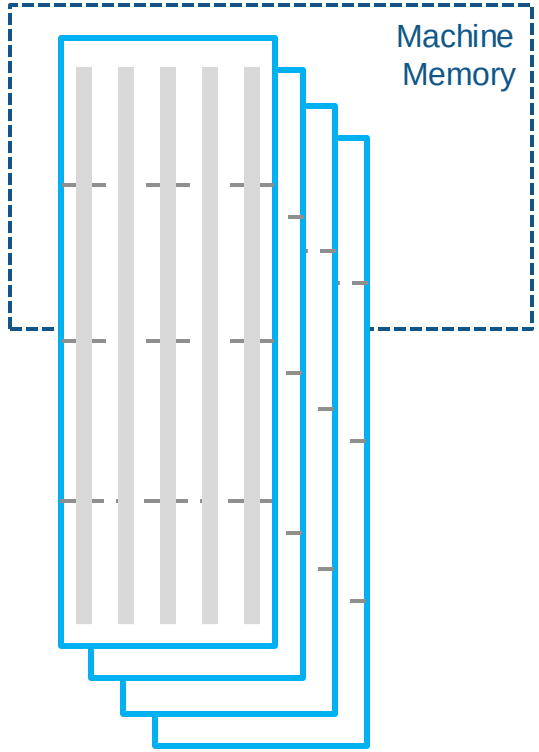

Tall Data

Execution Environments for Tall Arrays

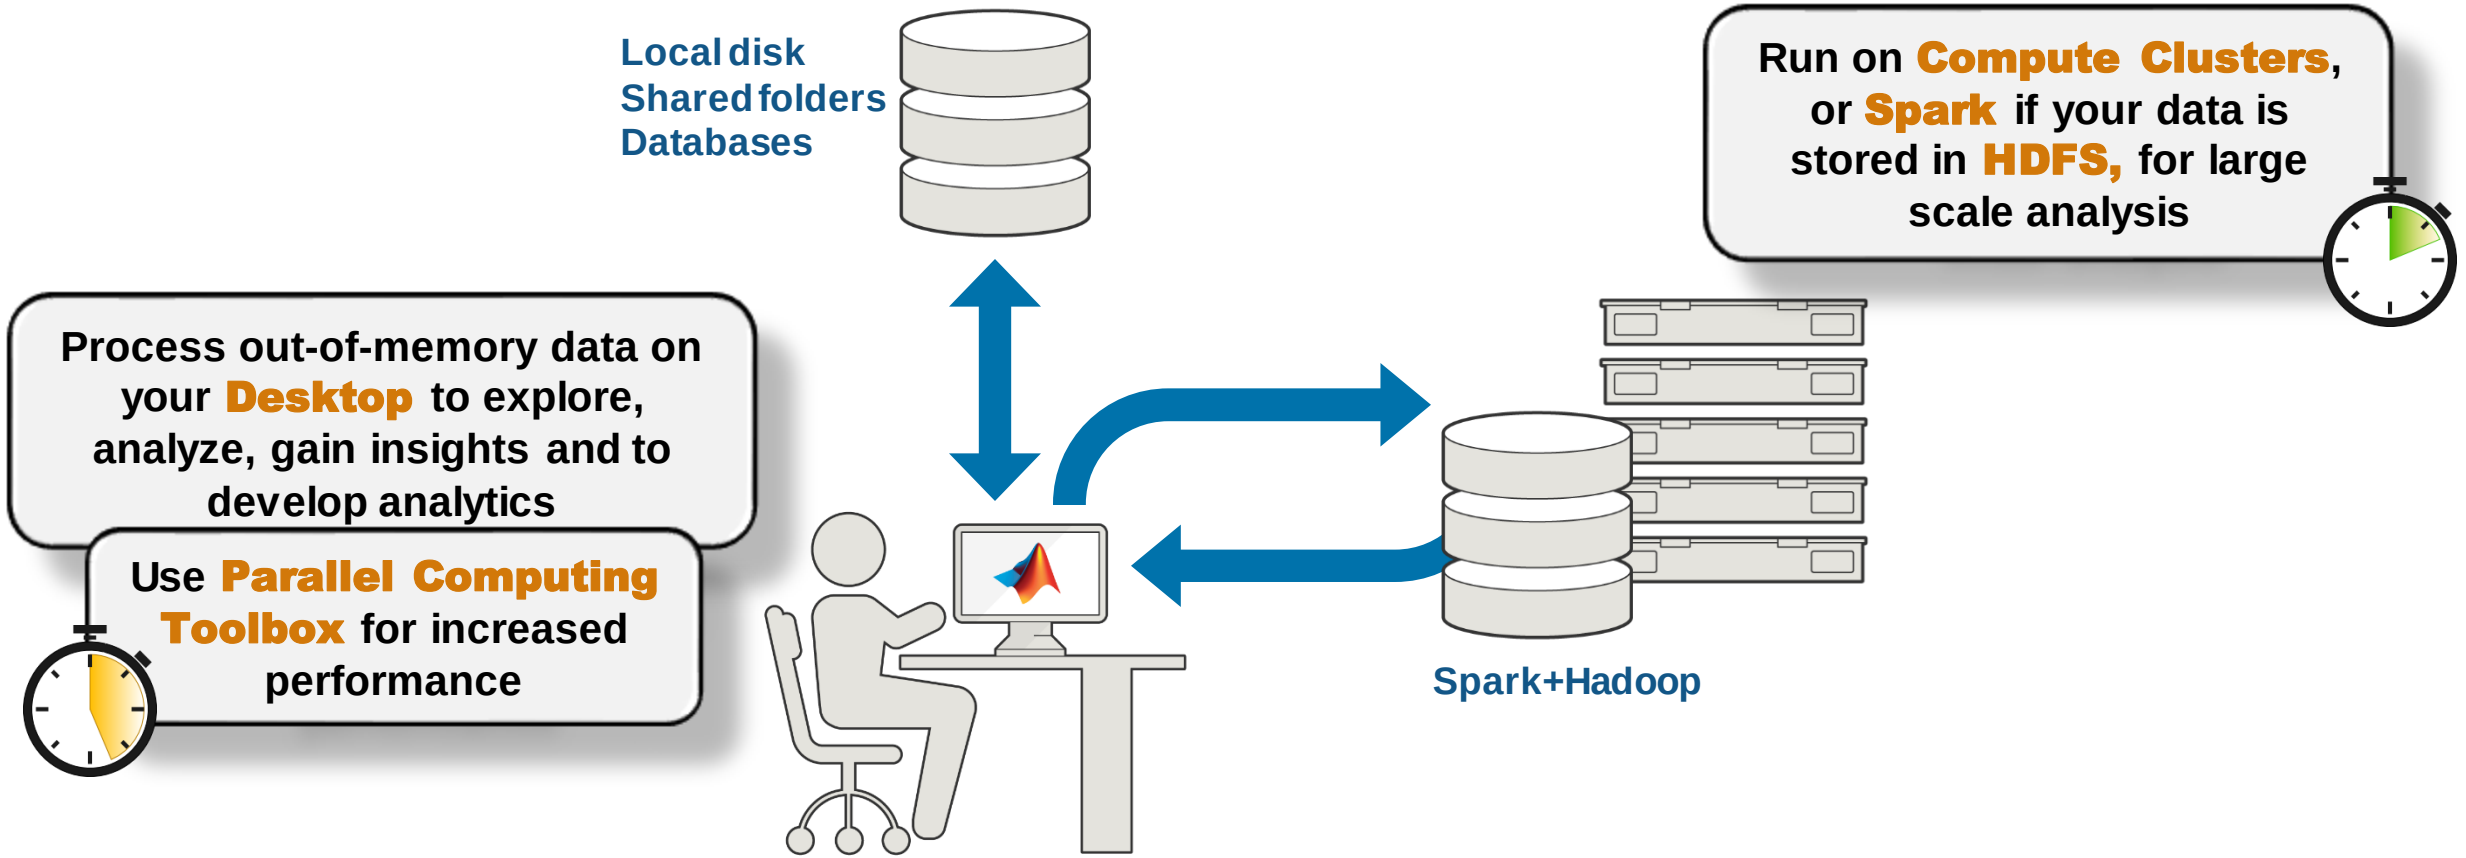

Distributed Arrays

- Distributed Arrays hold data remotely on workers running on a cluster
- Manipulate directly from client MATLAB (desktop)
- Common Actions
  - Matrix Manipulation
  - Linear Algebra and Signal Processing

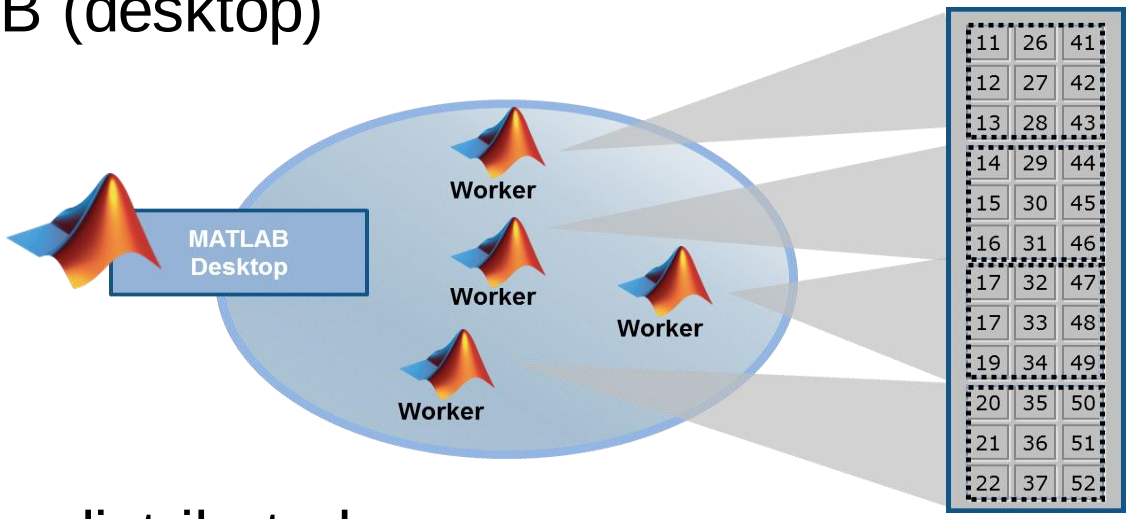

- 200+ MATLAB functions overloaded for distributed arrays

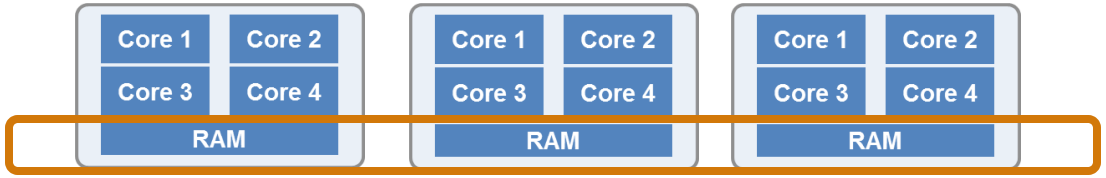

Supported functions for Distributed Arrays

Agenda

- Utilizing multiple cores on a desktop computer
- Scaling up to cluster and cloud resources
- Tackling data-intensive problems on desktops and clusters
- Accelerating applications with NVIDIA GPUs
- Summary and resources

Parallel Computing Paradigm

NVIDIA GPUs

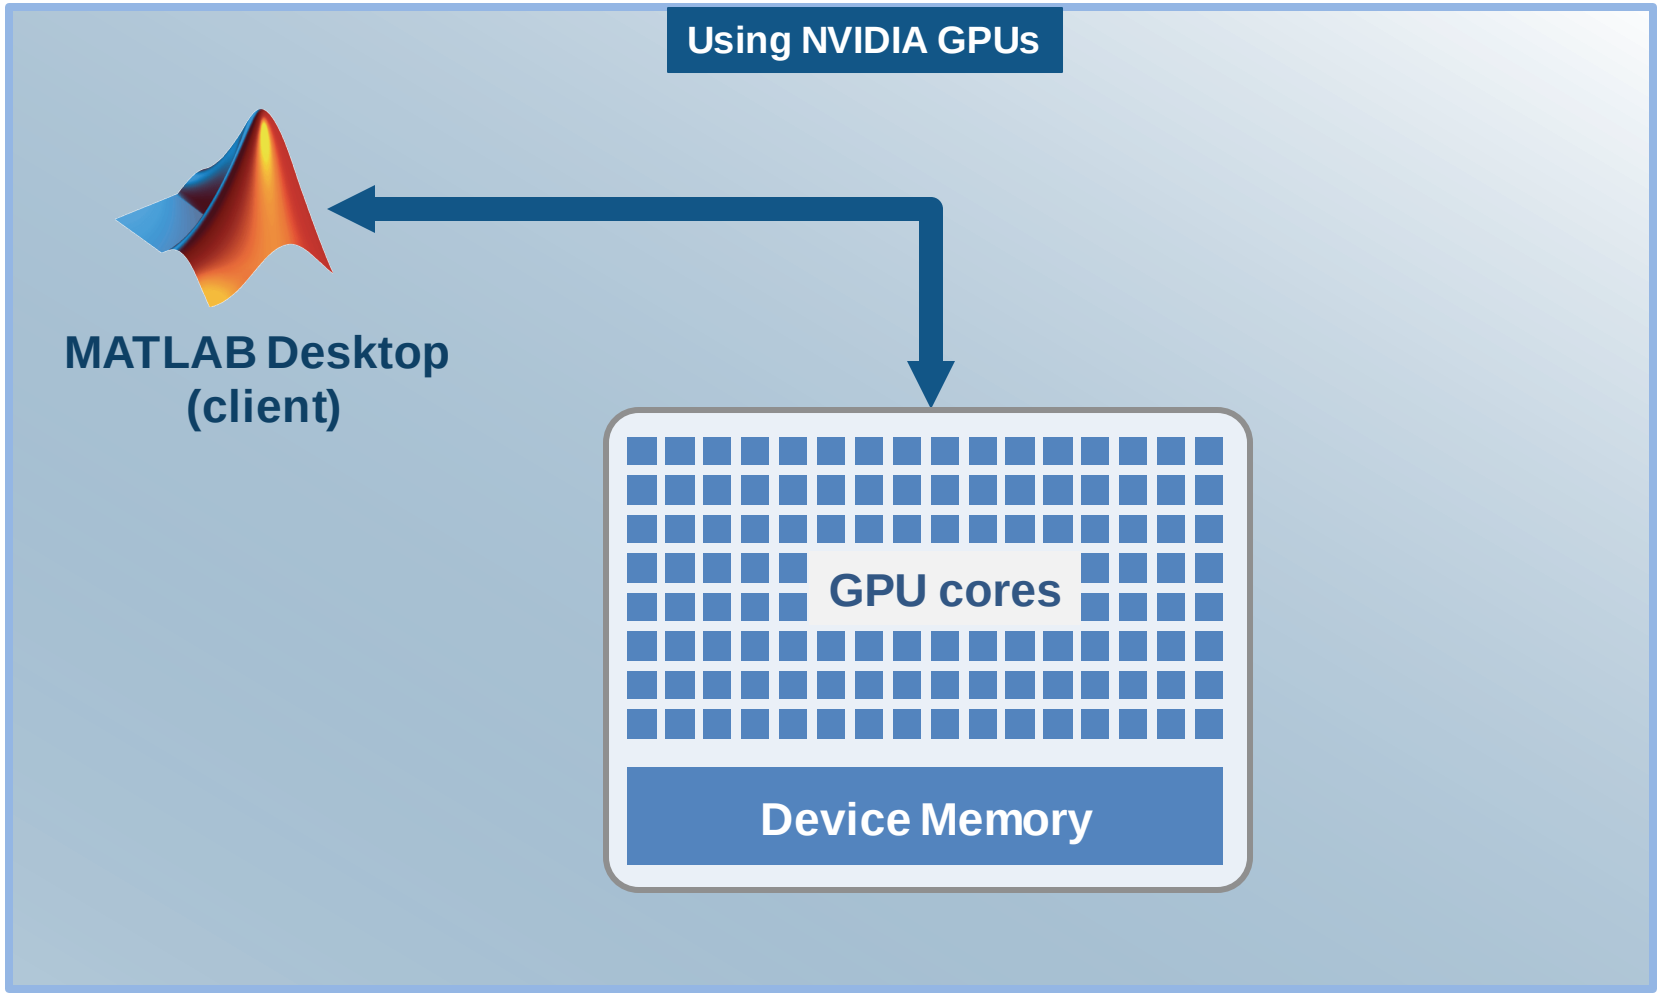

Programming with GPUs

- Built in toolbox support
- Simple programming constructs
  - `gpuArray, gather`

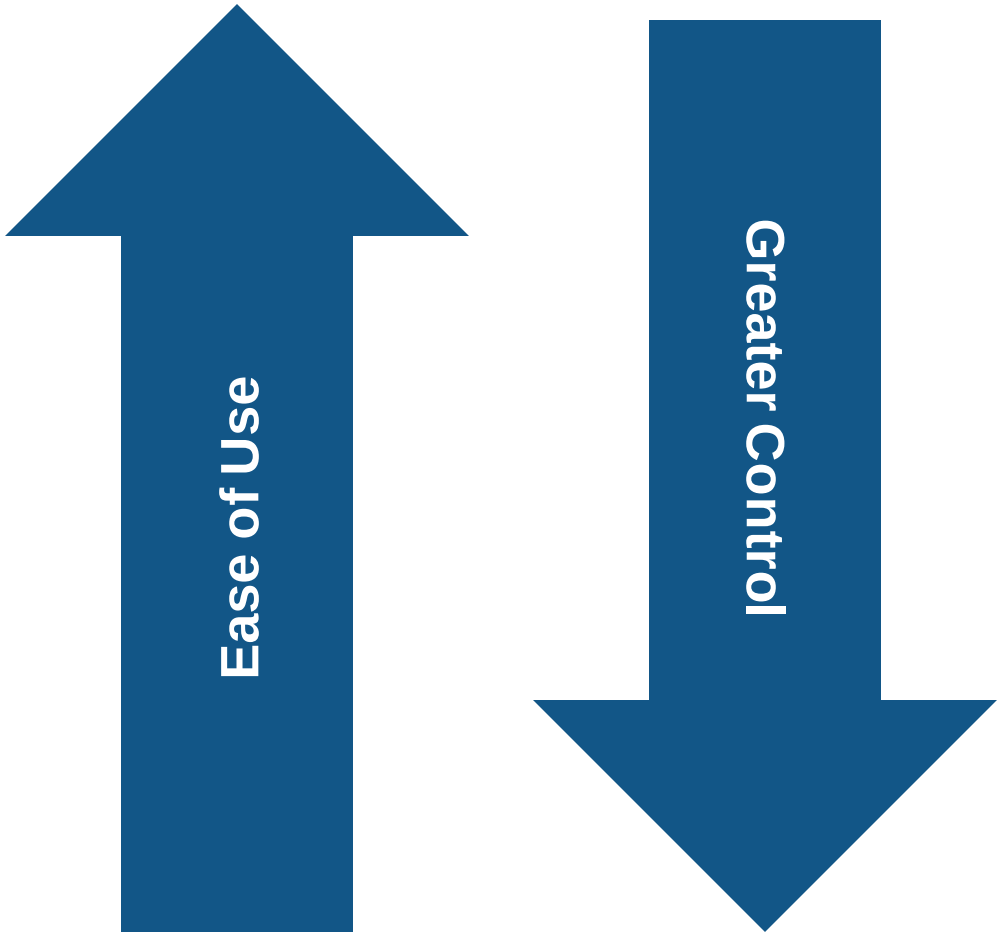

Programming with GPUs

- Built in toolbox support
- Simple programming constructs
  - `gpuArray, gather`
- Advanced programming constructs
  - `spmd, arrayfun`
- Interface for experts
  - `CUDAKernel, mex`

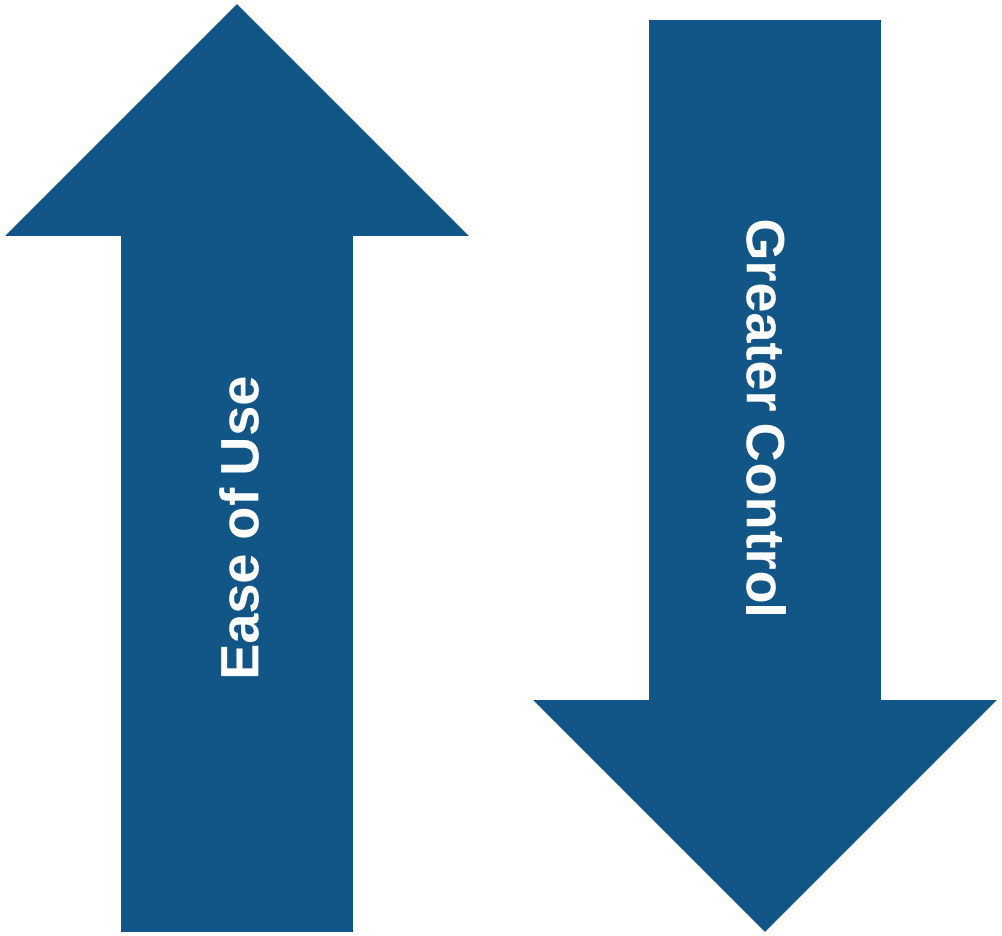

Speed-up using NVIDIA GPUs

- Ideal Problems
  - Massively Parallel and/or Vectorized operations
  - Computationally Intensive
  - Algorithm consists of supported functions
- 300+ GPU-enabled MATLAB functions
- Additional GPU-enabled Toolboxes
  - Algorithm consists of supported functions
  - Neural Networks
  - Image Processing
  - Communications
  - Signal Processing

  ..... Learn more

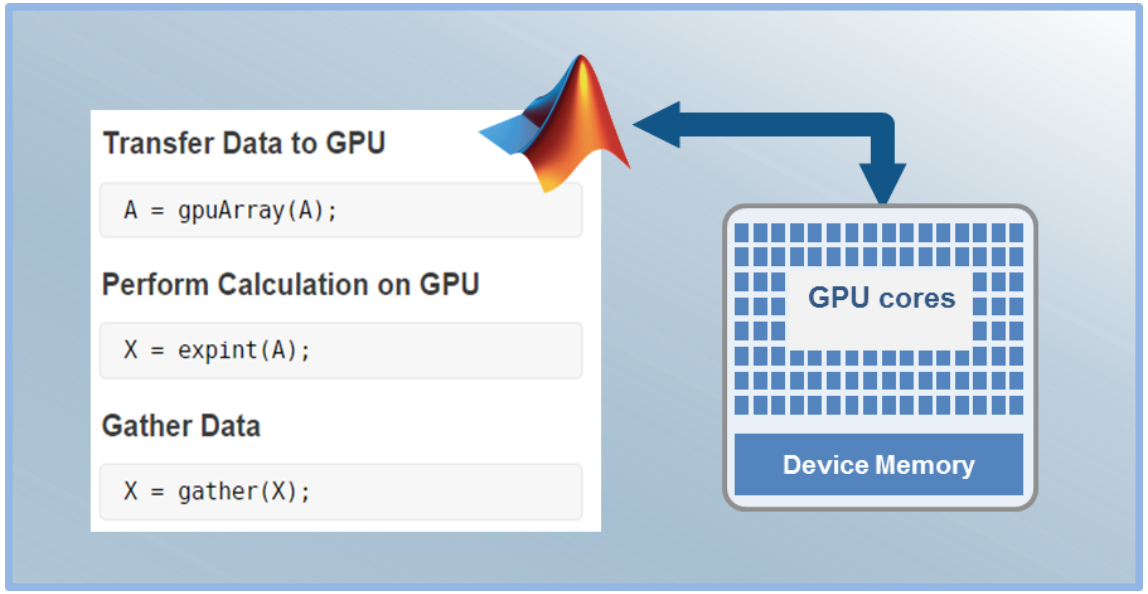

Agenda

- Utilizing multiple cores on a desktop computer
- Scaling up to cluster and cloud resources
- Tackling data-intensive problems on desktops and clusters
- Accelerating applications with NVIDIA GPUs
- Summary and resources

Summary

- Easily develop parallel MATLAB applications without being a parallel programming expert
- Speed up the execution of your MATLAB applications using additional hardware
- Develop parallel applications on your desktop and easily scale to a cluster when needed

Some Other Valuable Resources

- MATLAB Documentation
  - MATLAB → Advanced Software Development → Performance and Memory
  - Parallel Computing Toolbox
- Parallel and GPU Computing Tutorials
  - https://www.mathworks.com/videos/series/parallel-and-gpu-computing-tutorials-97719.html
- Parallel Computing on the Cloud with MATLAB
  - http://www.mathworks.com/products/parallel-computing/parallel-computing-on-the-cloud/

To Download the Workshop Files

- https://tinyurl.com/yb5brmuy

- Username: customer
- Password: MathWorks